

Machine Learning for Natural Language Processing

Neural Natural Language Processing

Lecture 4

Benjamin Muller

INRIA Paris - ALMANACH
benjamin.muller@inria.fr

- ① The Why and What of Natural Language Processing
- ② Representing text with vectors
- ③ Task specific Modeling of Text
- ④ Neural Natural Language Processing
- ⑤ Language Modeling
- ⑥ Transfer Learning with Neural Modeling for NLP

Task-specific modelling of textual data

- Preprocessing (encoding, segmentation)
- Modelling sequence tagging tasks (such as POS or NER) with MEMMS and CRF model
- Modelling Sentiment Analysis with a *bag-of-words* model
- Evaluating classification

- Deep Learning
 - Feed-Forward Neural Network
 - Recurrent Neural Network
 - Attention Mechanism
 - Embedding layer
- Training Neural Networks
- Differentiable Programming (Pytorch, Tensorflow)

Deep Learning toolkit

What is Deep Learning ?

Definition

*Deep learning is a class of machine learning algorithms that **uses multiple layers to progressively extract higher level features** from the raw input¹*

¹Deng et al. (2014)

This lecture is..

- This lecture **is not** a Deep Learning lecture
- This lecture is a Deep Learning lecture applied to Natural Language Processing

To grasp the overall Deep Learning picture Goodfellow et al. (2016)

Deep Learning and Natural Language Processing

- These last years, Deep Learning has become the main modelling framework for (almost) all NLP tasks

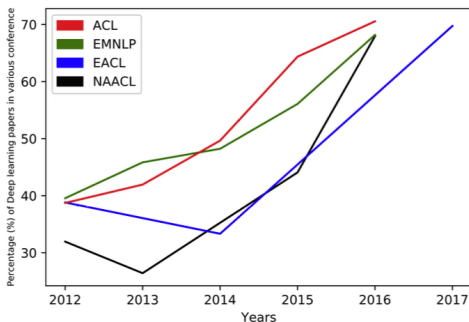


Figure: % of papers published in main NLP conferences that use deep learning

A Deep Learning Model is the combination of these three elements:

- Data
- Architecture
- Training Process (optimization)

- Logistic Regression (1 layer Feed-Forward Neural Network)
- Feed-Forward Neural Network
- Recurrent Neural Network
- Attention
- Embedding Layer

Let $(X_i, \dots, Y_i)_{i:1..n}$ $X_i \in R^d$, $Y_i \in R^{d'}$ input and output variables.

Goal: Learn a predictive model of Y_i with X_i .

Logistic Regression (1 layer Neural Network)

Assuming Y is single dimension ($d' = 1$)

$$Y = \sigma\left(\sum_j w_j X_j\right)$$

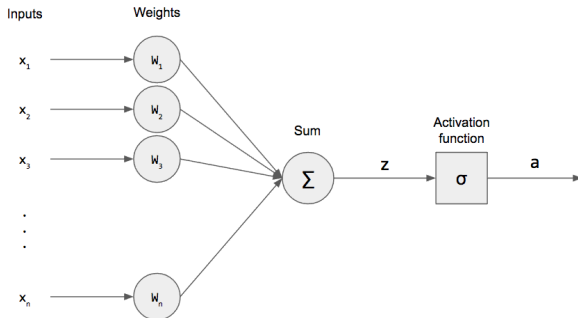


Figure: Representation of 1 layer Neural Network (NN) with sigmoid activation

Feed-Forward Neural Network²

We define a L layers network, with activation function f

W_l weights matrices $W_l \in R^{d_{l-1}, d_l}$

$$h_0 = x$$

$$h_l = f(W_l h_{l-1}) \text{ for } l \in 0, \dots, L$$

$$\hat{y} = h_L$$




Figure: Representation of 2 layers neural network (1 hidden layer)

NB : h_l is a **hidden vector**.

²also called **Multi-Layer Perceptron (MLP)**

Training Neural Networks: Backpropagation

- We train NNs with Stochastic Gradient Descent (SGD).
- For NNs, computing the gradient with the Chain Rule is very efficient computationally
- We call **backpropagation**, SGD applied to neural networks.
- All weights are **initialized randomly**

Training Neural Networks: Backpropagation

- We train NNs with Stochastic Gradient Descent (SGD).
- For NNs, computing the gradient with the Chain Rule is very efficient computationally
- We call **backpropagation**, SGD applied to neural networks.
- All weights are **initialized randomly**

Back-propagation for a Feed-Forward Neural Network (FNN)

Forward Pass for observation x_i

$$h_0 = x_i$$

$$h_l = f(W_l h_{l-1}) \text{ for } l \in 0, \dots, L$$

$$\hat{y}_i = h_L$$

$l_w(y_i, \hat{y}_i)$ loss function

Backward pass

$$l_w(y_i, \hat{y}_i)$$

$$w_{lk} := w_{lk} - \eta \frac{\partial l_w(y_i, \hat{y}_i)}{\partial w_{lk}}$$

$\frac{\partial l}{\partial w}$ computed with backprop.

Backpropagation

Backpropagation in a nutshell

- We alternate between forward pass i.e $x_i \rightarrow l_w(y_i, \hat{y}_i)$
- Backward pass during which each weight is updated with regard to the error $l_w(y_i, \hat{y}_i)$, $w \rightarrow w - \eta \frac{\partial l_w(y_i, \hat{y}_i)}{\partial w}$

Intuition

- Iteration after iteration: each weights are updated to minimize the loss function i.e the gap between the prediction and the observed values
- The network weights learn representation of the input that lead to the "best" prediction of the output.

Modelling Sequence with Neural Networks

Let $(X^1, \dots, X^T)_i, (Y^1, \dots, Y^T)_i$ sequence of input and output (e.g X^t *1-hot encoded words*, Y^t *POS tags*)

- Feed-Forward Neural Network do not model sequence as such
- Solution: Introducing *recurrent relation* in the network

- Family of neural network architectures in which a recurrent relation is induced by the architecture
- Recurrent Neural Network (RNN) Architectures
 - Vanilla RNN
 - Long-Short-Term Memory RNN

Vanilla RNN

Forward Pass for observation x_i ,
activation function f

$$h_0^t = x_i^t$$

for $l \in 0, \dots, L$ for $t \in 0, \dots, T$

$$h_l^t = f(W_l^f h_{l-1}^t + W_l^r h_l^{t-1} + b_l)$$

$$\hat{y}_i^t = h_L^t$$

$$l_w(y_i, \hat{y}_i) = \sum_{k=1..T} \sum_{k=1..d'} y_{ik}^t \cdot \log(\hat{y}_{ik}^t)$$




Figure: Recurrent Neural Network

Vanilla RNN schema




Figure: Recurrent Neural Network³

³<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- RNNs can be trained with some form of Backpropagation through time (BPTT)
- BPTT is the straight application of Backpropagation with recurrent connections⁴

⁴<https://machinelearningmastery.com/gentle-introduction-backpropagation-time/>

Limits of Vanilla RNN

- Vanishing Gradients problem.
First elements of the sequence will not get gradients updates. Hence the long range dependencies are poorly captured by the Vanilla RNN.
- Solution: **Memory specific weights** e.g Long-Short Term Memory (LSTM) cell

Let $x^t \in R^d$ element of the sequence (x^1, \dots, x^T) our input

We introduce *memory cell* variable C^t .

The recurrence formula of h^t based on h^{t-1} is given by :

$$\tilde{C}^t = \tanh(W_C[x_t, h_{t-1}] + b_c) \quad \text{candidate cell}$$

$$f^t = \sigma(W_f[x_t, h_{t-1}] + b_f) \quad \text{forget gate}$$

$$i^t = \sigma(W_i[x_t, h_{t-1}] + b_i) \quad \text{input gate}$$

$$o^t = \sigma(W_o[x_t, h_{t-1}] + b_o) \quad \text{ouput gate}$$

$$C^t = i^t \star \tilde{C}^t + f^t \star C^{t-1} \quad \text{new cell state}$$

$$h^t = o^t \star \tanh(C_t) \quad \text{new hidden vector}$$




Figure: LSTM ⁶

⁶<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- For each input , a new cell candidate is computed
- Based on the different gates, we compute what information is forgotten, took from the new candidate, and outputed.
- Every weights are trained simultaneously with BPTT
- Based on the input - output, every gates and weights are updated to lead to the "best" prediction

- Each hidden vector h^t captures information about the left sequence
- Each hidden vector represents the given token t based on its left context.
- h^t is a **contextual** representation of the token t

- Feed-Forward Layers, Recurrent Layers (Vanilla, LSTMs) can be used as building blocks for any more complex architectures
- Architectural decisions are task specific
- Best practices and Empirical results should drive what architecture is chosen for a given task

How to model discrete data with Neural Networks ?

- So far, we did not focus on the input data
- Neural Networks work better have been shown to perform better with dense vectors compared discrete vectors
- How to handle discrete tokens tokens efficiently ?
→ embedding layer

Embedding Layer

- A **token** is the **basic unit** of discrete data, defined to be an item from a vocabulary indexed by $1, \dots, V$.
- We define $E \in R^{Vd}$ an embedding matrix which associates each token to a vector in R^d
- We train those weights simultaneously with backpropagation as any other weights
- We initialize this embedding layer:

Embedding Layer

- A **token** is the **basic unit** of discrete data, defined to be an item from a vocabulary indexed by $1, \dots, V$.
- We define $E \in R^{V \times d}$ an embedding matrix which associates each token to a vector in R^d
- We train those weights simultaneously with backpropagation as any other weights
- We initialize this embedding layer:
 - randomly as any other weights
 - using a pretrained skip-gram word2vec (or any other word embedding techniques)

Sequence Labelling with LSTM: back to POS tagging

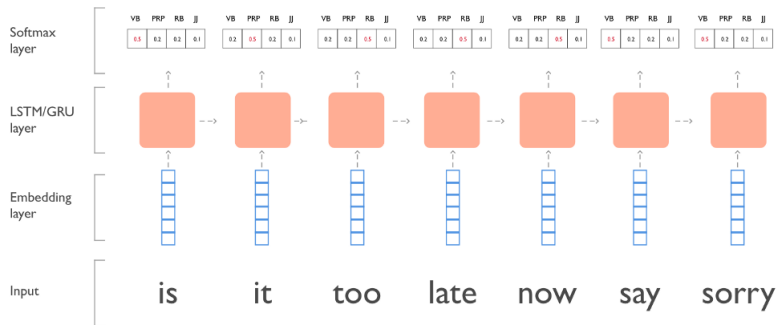


Figure: POS tagging with embedding layer and LSTM recurrent network

7

⁷GRU is a light variant of LSTM

Sequence Labelling with LSTM: back to POS tagging

We define a sequence labelling RNN with LSTM cell and output feed-forward layer

Let (x^0, \dots, x^T) sequence of words, (y^0, \dots, y^T) tags.

All LSTM weights, W_{fo} forward layer and Embedding layer E initialized randomly

For $t \in 0, \dots, T$

$$h^t, C^t = LSTM_w(x^t, h^{t-1}, C^{t-1})$$

$$s^t = \tanh(W_{fo} h^t)$$

$$\hat{y}^t = \operatorname{argmax}_k(\hat{p}^t)$$

We compute the Cross-Entropy loss based on

$$l_w(y_i, \hat{y}_i) = \sum_{k=1..T} \sum_{k=1..d'} y_{ik}^t \cdot \log(\hat{p}_{ik}^t)$$

We apply backpropagation.

- RNNs define a recurrence relation in the left-to-right direction
- What if we need both left-to-right and right-to-left context
→ **bi-directional** LSTM

Bi-Directional LSTMs




Figure: LSTM

- Define two LSTM cells: one left to right and one right to left
- Concatenate for each step the output hidden vectors

- For input element x^t : the RNN provides a vector representation h^t (contextual)
- For some tasks, some input are more important than others
→ attention mechanism

Let (x^0, \dots, x^T) sequence of words, y label (sequence classification)

We define LSTM+dense layer neural network

$$s^T = f_W(h^T) \text{ with } h^t, C^t = LSTM(h^{t-1}, C^{t-1}, x^t) \text{ for all } t$$

With Attention it becomes

$$s^T = f_W(\widetilde{h^T}) \text{ with } \widetilde{h^T} = \sum_t a_t h^t \text{ and } h^t = LSTM(h^{t-1}, C^{t-1}, x^t)$$

Question: How to compute a_t ?

How to compute Attention weights ?

Here is a simple attention mechanism:

$$u^t = \tanh(W_a h^t)$$

$$a^t = \text{softmax}_w(u^t)$$

$$\widetilde{h^T} = \sum_t a_t h^t$$

- Many different variants of attention mechanisms
- As with all other weights, attention weights are trained end-to-end with backpropagation
- \widetilde{h}^T is also a contextual representation vector of the token x^T

Attention mechanism as an interpretation tool

- Neural networks are black box models
- Attention mechanisms provide some insights on what information the network is using to make its prediction
- Warning: It does not say how it is using it!

Attention mechanism as an interpretation tool

Using the attention weight we can compute heatmaps in a straightforward way :




Figure 3: Attention visualization for a review sentence

Zhou et al. (2016)

How do we build/train neural networks in practice ?

- All seen neural networks can be trained using backpropagation ⁸
- Implementing backpropagation by hand for all our models is time consuming

⁸if activation function are differential

How do we build/train neural networks in practice ?

- All seen neural networks can be trained using backpropagation ⁸
- Implementing backpropagation by hand for all our models is time consuming
- Happily, came **differentiable programming** and deep learning libraries!

⁸if activation function are differential

Definition

Differentiable programming is a programming paradigm in which the programs can be differentiated throughout, usually via automatic differentiation

Deep Learning libraries are implementation of automatic differentiation modules optimized with user-friendly tools for Deep Learning.

e.g: Tensorflow, Keras, Dynet, Pytorch are popular Deep Learning libraries In the lab we will use **Pytorch**

With deep learning libraries, implementing neural networks becomes easy

- Pre-implemented modules like: *Linear* layers, LSTM cells, SGD optimizers
- Automatic Differentiation

In the following lab, we will build a LSTM recurrent neural network for Sentiment Classification using pytorch.

- Deep Learning architectures for NLP
 - Feed-Forward Neural Network
 - Recurrent Neural Network (LSTM)
 - Embedding layer
 - Attention Mechanism
- Deep Learning libraries

- Deng, L., Yu, D., et al. (2014). Deep learning: methods and applications. *Foundations and Trends® in Signal Processing*, 7(3–4):197–387.
- Goodfellow, I., Bengio, Y., and Courville, A. (2016). *Deep Learning*. MIT Press. <http://www.deeplearningbook.org>.
- Zhou, X., Wan, X., and Xiao, J. (2016). Attention-based lstm network for cross-lingual sentiment classification. In *Proceedings of the 2016 conference on empirical methods in natural language processing*, pages 247–256.