

Machine Learning for Natural Language Processing

Sequence Generation

Lecture 5

Benjamin Muller

INRIA Paris - ALMANACH
benjamin.muller@inria.fr

- ① The Why and What of Natural Language Processing
- ② Representing text with vectors
- ③ Task specific Modeling of Text
- ④ Neural Natural Language Processing
- ⑤ Language Modelling
- ⑥ Transfer Learning

- Language Model
- Conditioned Language Model : focus on Sequence to Sequence

Language Model

- What is a Language Model ?
- Modelling language with n-grams
- Modelling language with a LSTM
- The Transformer Architecture

Language modeling

What is language modeling?

- **Language modeling** corresponds to assigning a probability to a text
- A text is a **sequence of tokens**, or characters
- Tokens can be words, sub-words,
- For example:

$$\{\text{a cat}\} = \{\text{a, cat}\},$$

What is language modeling?

- **Language modeling** corresponds to assigning a probability to a text
- A text is a **sequence of tokens**, or characters
- Tokens can be words, sub-words,
- For example:

$$\begin{aligned}\{\text{a cat}\} &= \{\text{a, cat}\}, \\ &= \{\text{a, }, \text{c, a, t}\},\end{aligned}$$

What is language modeling?

- **Language modeling** corresponds to assigning a probability to a text
- A text is a **sequence of tokens**, or characters
- Tokens can be words, sub-words,
- For example:

$$\begin{aligned}\{\text{a cat}\} &= \{\text{a, cat}\}, \\ &= \{\text{a, }, \text{c, a, t}\}, \\ &= \{\text{a, }, \text{ca, t}\}.\end{aligned}$$

What is language modeling?

- Given a sequence $\{w_1, \dots, w_T\}$ of tokens, a language model estimates its probability:

$$P(w_1, \dots, w_T)$$

- P depends on a **vocabulary**, i.e., the set of unique tokens.
- Question: How to estimate P ?

What is language modeling?

- Given a sequence $\{w_1, \dots, w_T\}$ of tokens, a language model estimates its probability:

$$P(w_1, \dots, w_T)$$

- P depends on a **vocabulary**, i.e., the set of unique tokens.
- Question: How to estimate P ?

Language Models

- Causal Language Model
- Mask Language Model

Language models are applied in several fields:

- Speech recognition:

$$P(\text{"Vanilla, I scream"}) < P(\text{"Vanilla ice cream"}).$$

- Machine translation:

$$P(\text{"D   u en bien"} \mid \text{"Pleasantly surprised"}) < \\ P(\text{"Agr  ablement surpris"} \mid \text{"Pleasantly surprised"})$$

- Optical Character Recognition:

$$P(\text{"m0ve fast"}) < P(\text{"move fast"})$$

Probabilistic *Causal* language model

- Sequence probability as a product of token probabilities:

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{t-1}, \dots, w_1)$$

Probabilistic *Causal* language model

- Sequence probability as a product of token probabilities:

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{t-1}, \dots, w_1)$$

- Indeed we have:

$$P(a, b) = P(a)P(b \mid a)$$

Probabilistic *Causal* language model

- Sequence probability as a product of token probabilities:

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{t-1}, \dots, w_1)$$

- Indeed we have:

$$P(a, b) = P(a)P(b \mid a)$$

- Recursively applied to a sequence:

$$\begin{aligned} P(w_1, w_2, w_3) &= P(w_1)P(w_2, w_3 \mid w_1) \\ &= P(w_1)P(w_2 \mid w_1)P(w_3 \mid w_2, w_1). \end{aligned}$$

Probabilistic *Causal* language model

- Sequence probability as a product of token probabilities:

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{t-1}, \dots, w_1)$$

- Indeed we have:

$$P(a, b) = P(a)P(b \mid a)$$

- Recursively applied to a sequence:

$$\begin{aligned} P(w_1, w_2, w_3) &= P(w_1)P(w_2, w_3 \mid w_1) \\ &= P(w_1)P(w_2 \mid w_1)P(w_3 \mid w_2, w_1). \end{aligned}$$

- Causal Language models estimate probability of upcoming token given past:

$$P(w_t \mid w_{t-1}, \dots, w_1).$$

Estimating Language Models

- Causal Language Model
- Mask Language Model

Sentence The cat is drinking milk in the kitchen

¹ Devlin et al. (2018)

²also referred as Cloze Task

Sentence The cat is drinking milk in the kitchen

input The cat <MASK> drinking <MASK> in the kitchen

- Randomly replace 15% of words in sentence with a <MASK> token

¹ Devlin et al. (2018)

²also referred as Cloze Task

Sentence The cat is drinking milk in the kitchen

input The cat <MASK> drinking <MASK> in the kitchen

targets {"is", "milk"}

- Randomly replace 15% of words in sentence with a <MASK> token
- Take the masked words as targets for the model to predict

¹ Devlin et al. (2018)

²also referred as Cloze Task

Sentence The cat is drinking milk in the kitchen

input The cat **mushroom** drinking **shoes** in the kitchen

targets {"is", "milk"}

- Randomly replace 15% of words in sentence with a <MASK> token
- Take the masked words as targets for the model to predict
- Extension: use random words from vocabulary instead of <MASK>

¹ Devlin et al. (2018)

²also referred as Cloze Task

Masked Language Modelling estimates the probability of sequence tokens of length T with:

$$P(w_i | w_1, \dots, w_{i-1}, w_i, \dots, w_T)$$

Language Models in a nutshell

- a Language Model is a model that predicts a **token** based on its surrounding linguistic **context**
- Tokens can be words, sub-words or characters
- Context can be the *left sequence*, *left and right sequence*, the sentence, a window around the words, the paragraph...
- We saw two way of defining language models: Causal Language Model and Mask Language Model

Estimating language models

- Statistical approach: N-Gram model
- Neural Language Models
 - Recurrent Neural Networks (LSTM)
 - The Transformer Architecture

- Example:

$$\begin{aligned}P(\text{English} \mid \text{The moment one learns}) &= \frac{c(\text{The moment one learns English})}{c(\text{The moment one learns})} \\&= \frac{35}{73} = 0.48\end{aligned}$$

Sentence “The moment one learns English” appears 35 in dataset

Sentence “The moment one learns” appears 75 in dataset

Limitations of count based language model

- Number of unique sentences increases with dataset size,
 - Long sentences are rare: no good statistics for them
- Too many sentences with not enough statistics
(Sparsity due to combinatorial structure of language)

Count based language model

- **Solution** truncate past to a fixed size window
- For example:

$$P(\text{English} \mid \text{The moment one learns}) \approx P(\text{English} \mid \text{one learns})$$

- Implicit assumption: **most important information about a word is in its recent history**
- **Beware!** In general:

$$P(w_1, \dots, w_T) \neq \prod_{t=1}^T P(w_t \mid w_{t-1}, \dots, w_{t-n+1})$$

- **Truncated count based models = n -gram models**
- “ n ” refers to the size of past
- Examples:
 - Unigram:

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t)$$

- Bigram:

$$P(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{t-1})$$

Count based language model: unigram

- Probability of a sentence with a unigram model:

$$P_U(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t) = \prod_{t=1}^T \frac{c(w_t)}{N}$$

N = total number of tokens in dataset

$c(w_t)$ = number of occurrences of w_t in dataset

- Unigram only uses word frequency
- Example of text generation with this model:

the or is ball then car

Count based language model: bigram

- Probability of a sentence with a bigram model:

$$P_U(w_1, \dots, w_T) = \prod_{t=1}^T P(w_t \mid w_{t-1}) = \prod_{t=1}^T \frac{c(w_{t-1}w_t)}{c(w_{t-1})}$$

$c(w_{t-1}w_t)$ = number of occurrences of sequence $w_{t-1}w_t$

- Predict a word just with the previous word

Count based language model: bigram

- Example of text generation with bigram model:

new car parking lot of the

- “car” is generated from “new”, “parking” from “car”...
- But “new” has no influence on “parking”

Count based language model

- Simple to extend to longer dependencies: trigrams, 4-grams...
- n -grams can be “good enough” in some cases
- But n -grams cannot capture long term dependencies required to truly model language

Estimating n -gram probabilities: an example

- bigram:

$$P(w_t \mid w_{t-1}) = \frac{c(w_{t-1}w_t)}{c(w_{t-1})}$$

- Dataset:

<s>we sat in the house

<s>we sat here we two and we said

<s>how we wish we had something to do

- Extract some probabilities:

$$P(\text{sat} \mid \text{we}) = 0.33, \quad P(\text{wish} \mid \text{we}) = 0.17, \quad P(\text{in} \mid \text{sat}) = 0.5$$

- <s> = token for beginning of sentence; $P(\text{<s>}) = 1$.
- Compute sentence probability with them

Estimating n -gram probabilities: an example

- Extract count from Berkeley Restaurant dataset (9222 sentences)
- Unigram counts:

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

- Bigram counts:

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

Estimating n -gram probabilities: an example

- The bigram probabilities are obtained by dividing the bigram counts with the unigram counts:

$$P(w_2 \mid w_1) = \frac{c(w_1 w_2)}{c(w_1)}$$

- Resulting bigram probabilities:

	i	want	to	eat	chinese	food	lunch	spend
i	0.022	0.33	0	0.036	0	0	0	0.00079
want	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011
to	0.00083	0	0.0017	0.28	0.00083	0	0.0025	0.087
eat	0	0	0.0027	0	0.021	0.0027	0.056	0
chinese	0.0063	0	0	0	0	0.52	0.0063	0
food	0.014	0	0.014	0	0.00092	0.0037	0	0
lunch	0.0059	0	0	0	0	0.0029	0	0
spend	0.0036	0	0.0036	0	0	0	0	0

Estimating n -gram probabilities: an example

- Example:

$$P(< s> \text{ i want chinese food})?$$

$< s>$ = token for beginning of sentence; $P(< s>) = 1$.

- Result:

$$\begin{aligned} P(< s> \text{ i want chinese food}) &= P(< s>)P(\text{i} | < s>)P(\text{want} | \text{i})P(\text{chinese} | \text{want})P(\text{food} | \text{chinese}) \\ &= 1 \times .25 \times 0.33 \times 0.0065 \times 0.52 \\ &= 0.00027885 \end{aligned}$$

Estimating n -gram probabilities: an example

	i	want	to	eat	chinese	food	lunch	spend
i	0.022	0.33	0	0.036	0	0	0	0.00079
...								
food	0.014	0	0.014	0	0.00092	0.0037	0	0
lunch	0.0059	0	0	0	0	0.0029	0	0
spend	0.0036	0	0.0036	0	0	0	0	0

- Example:

$$P(<s> \text{ i bring my lunch to work})?$$

- Result:

$$\begin{aligned}P(<s> \text{ i bring my lunch to work}) &= P(<s>) \dots P(\text{to}|\text{lunch}) \dots \\&= 1 \times \dots \times 0 \times \dots \\&= \mathbf{0}\end{aligned}$$

- **Does not generalize well!**

Good-Turing estimation

- **Idea** reallocate probability mass of n -grams that occur exactly $c + 1$ times to n -grams that occur exactly c times
 - reallocate mass of n -grams appearing once to unseen n -grams
- **alternative to Add-1**

- **Idea** reallocate probability mass of n -grams that occur exactly $c + 1$ times to n -grams that occur exactly c times
- reallocate mass of n -grams appearing once to unseen n -grams

→ **alternative to Add-1**

- the adjusted count:

$$c^* = (c + 1) \frac{N_{c+1}}{N_c}$$

where N_c is the number of n -grams that appears exactly c times

Good-Turing estimation

- **Idea** reallocate probability mass of n -grams that occur exactly $c + 1$ times to n -grams that occur exactly c times
- reallocate mass of n -grams appearing once to unseen n -grams

→ **alternative to Add-1**

- the adjusted count:

$$c^* = (c + 1) \frac{N_{c+1}}{N_c}$$

where N_c is the number of n -grams that appears exactly c times

- n -gram probability depends on c^* instead of c

- **Idea** reallocate probability mass of n -grams that occur exactly $c + 1$ times to n -grams that occur exactly c times
- reallocate mass of n -grams appearing once to unseen n -grams

→ **alternative to Add-1**

- the adjusted count:

$$c^* = (c + 1) \frac{N_{c+1}}{N_c}$$

where N_c is the number of n -grams that appears exactly c times

- n -gram probability depends on c^* instead of c
- **Problem** What if $N_{c+1} = 0$ (but $N_c > 0$)?

- If no good statistics on long context: use shorter context
- **Backoff**: use trigram if enough data, else backoff to bigram.
- **Interpolation**: mix statistics of trigram, bigram and unigram.

Pros and Cons of N-Gram Language Models

Pros

- Fast at training and inference
- Can reach good accuracy if lots of data

Cons

- Impossible to model very long dependencies (simplistic assumptions done)
- Generalization limited
- Not Deep Learning compatible

- Statistical approach: N-Gram model
- Neural Language Models
 - Recurrent Neural Networks (LSTM)
 - The Transformer Architecture

How to frame language modeling in a deep learning compatible way ?
What neural architecture/objective ?

- Neural Language Model objective and Training
- Architectures
 - Recurrent Network
 - Transformer

Neural Language Model

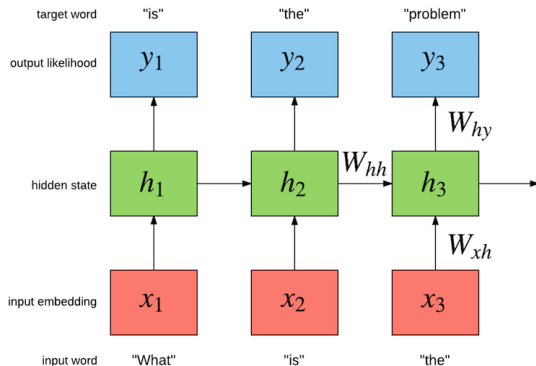


Figure: Neural Language Modelling schema view ³

³<http://torch.ch/blog/2016/07/25/nce.html>

Neural Language Model training and inference

Let (x^1, \dots, x^T) ; sequence of tokens (1-hot encoded), E embedding layer,
 NN_θ a sequential model (e.g. LSTM) $\{f, W\}$ dense layer

We present forward/backward step to predict token x^{t+1} with x^1, \dots, x^t

$$e_t = Ex^t \quad \forall t \leq T \quad \text{Embedding layer}$$

$$h_t = NN_\theta(e_1, \dots, e_{t-1}, e_t) \quad \text{sequential layers with weights } \theta$$

$$s_t = f(Wh_t) \quad \text{Dense Layer}$$

Neural Language Model training and inference

Let (x^1, \dots, x^T) ; sequence of tokens (1-hot encoded), E embedding layer,
 NN_θ a sequential model (e.g. LSTM) $\{f, W\}$ dense layer

We present forward/backward step to predict token x^{t+1} with x^1, \dots, x^t

$$e_t = Ex^t \quad \forall t \leq T \quad \text{Embedding layer}$$

$$h_t = NN_\theta(e_1, \dots, e_{t-1}, e_t) \quad \text{sequential layers with weights } \theta$$

$$s_t = f(Wh_t) \quad \text{Dense Layer}$$

Train time

Neural Language Model training and inference

Let (x^1, \dots, x^T) ; sequence of tokens (1-hot encoded), E embedding layer,
 NN_θ a sequential model (e.g. LSTM) $\{f, W\}$ dense layer

We present forward/backward step to predict token x^{t+1} with x^1, \dots, x^t

$$e_t = Ex^t \quad \forall t \leq T \quad \text{Embedding layer}$$

$$h_t = NN_\theta(e_1, \dots, e_{t-1}, e_t) \quad \text{sequential layers with weights } \theta$$

$$s_t = f(Wh_t) \quad \text{Dense Layer}$$

Train time

$$\hat{p}_t = \text{softmax}_V(o_t) = \left(\frac{e^{o_{tv}}}{\sum_k e^{o_{tk}}} \right)_{v \in 1..V}$$

Neural Language Model training and inference

Let (x^1, \dots, x^T) ; sequence of tokens (1-hot encoded), E embedding layer,
 NN_θ a sequential model (e.g. LSTM) $\{f, W\}$ dense layer

We present forward/backward step to predict token x^{t+1} with x^1, \dots, x^t

$$e_t = Ex^t \quad \forall t \leq T \quad \text{Embedding layer}$$

$$h_t = NN_\theta(e_1, \dots, e_{t-1}, e_t) \quad \text{sequential layers with weights } \theta$$

$$s_t = f(Wh_t) \quad \text{Dense Layer}$$

Train time

$$\hat{p}_t = \text{softmax}_V(o_t) = \left(\frac{e^{o_{tv}}}{\sum_k e^{o_{tk}}} \right)_{v \in 1..V}$$

$$\text{loss} = CE(\hat{p}_t, p_t) = -\log(\hat{p}_{t_{x^{t+1}}})$$

Neural Language Model training and inference

Let $(x^1, \dots, x^T)_i$ sequence of tokens (1-hot encoded), E embedding layer,
 NN_θ a sequential model (e.g. LSTM) $\{f, W\}$ dense layer

We present forward/backward step to predict token x^{t+1} with x^1, \dots, x^t

$$e_t = Ex^t \quad \forall t \leq T \quad \text{Embedding layer}$$

$$h_t = NN_\theta(e_1, \dots, e_{t-1}, e_t) \quad \text{sequential layers with weights } \theta$$

$$s_t = f(Wh_t) \quad \text{Dense Layer}$$

Train time

$$\hat{p}_t = \text{softmax}_V(o_t) = \left(\frac{e^{o_{tv}}}{\sum_k e^{o_{tk}}} \right)_{v \in 1..V}$$

$$\text{loss} = CE(\hat{p}_t, p_t) = -\log(\hat{p}_{t_{x^{t+1}}})$$

Compute ∇loss backprop
(update E, θ, W)

Neural Language Model training and inference

Let $(x^1, \dots, x^T)_i$ sequence of tokens (1-hot encoded), E embedding layer,
 NN_θ a sequential model (e.g. LSTM) $\{f, W\}$ dense layer

We present forward/backward step to predict token x^{t+1} with x^1, \dots, x^t

$$e_t = Ex^t \quad \forall t \leq T \quad \text{Embedding layer}$$

$$h_t = NN_\theta(e_1, \dots, e_{t-1}, e_t) \quad \text{sequential layers with weights } \theta$$

$$s_t = f(Wh_t) \quad \text{Dense Layer}$$

Train time

$$\hat{p}_t = \text{softmax}_V(o_t) = \left(\frac{e^{o_{tv}}}{\sum_k e^{o_{tk}}} \right)_{v \in 1..V}$$

Inference/Prediction Time

$$x_{t+1} = \text{argmax}_{v \in 1, \dots, V}(s_{tv})$$

$$\text{loss} = CE(\hat{p}_t, p_t) = \log(\hat{p}_{t_{x^{t+1}}})$$

Compute ∇loss backprop
(update E, θ, W)

Neural Language Model with LSTM cell

In this case, NN_θ is defined as (seen in lecture 4):
(θ is equal to $W_{p \in C, f, i, o}$)

Based on e_1, \dots, e_t we compute iteratively h_1, \dots, h_t

$$\tilde{C}^t = \tanh(W_C[e_t, h_{t-1}] + b_c) \quad \text{candidate cell}$$

$$f^t = \sigma(W_f[x_t, h_{t-1}] + b_f) \quad \text{forget gate}$$

$$i^t = \sigma(W_i[x_t, h_{t-1}] + b_i) \quad \text{input gate}$$

$$o^t = \sigma(W_o[x_t, h_{t-1}] + b_o) \quad \text{output gate}$$

$$C^t = i^t \star \tilde{C}^t + f^t \star C^{t-1} \quad \text{new cell state}$$

$$h_t = o^t \star \tanh(C_t) \quad \text{new hidden vector}$$

- Language Models are evaluated with perplexity

$$\textit{perplexity} = 2^{-p_i \log(\hat{p}_i)}$$

- It is a measure of "surprise" of the model

Comparing various language Models

Model	Perplexity
Kneser-Ney 5-gram	141
Neural n -gram	140
RNN	125
LSTM	115

- Penn Treebank dataset
- LSTM outperforms RNN

Limits of LSTM-based architectures

- LSTM models are widely used in NLP for their ability to model sequential data
- In theory, they are able to model sequences of infinite length (Siegelmann and Sontag, 1992)
- In practice, until recently LSTM based models were State-of-the-Art (SOTA) for language modeling (Rae et al., 2018)
- In practice, the recurrent nature of LSTM limits the possibility to scale the training process to more data (we cannot parallelize LSTM easily!)
- → **Transformer** were recently shown to work better in practice for a great variety of tasks including Language Model (Radford et al., 2019)

The Transformer Architecture⁵

⁵Vaswani et al. (2017)

Combining vectors with attention

- Use (self) attention mechanism
- Given a set of vectors $\mathbf{w}_1, \dots, \mathbf{w}_T \in \mathbb{R}^d$ representing words

$$\mathbf{h}_t = \sum_{i=1}^T a_{it} \mathbf{V} \mathbf{w}_i$$

where $\sum_{i=1}^T a_{it} = 1$.

- We could use $a_{it} = \frac{1}{T}$ and get bag of words
- We can also learn a_{it} based on the input and output as we did for the standard attention mechanism

- Introducing matrix $\mathbf{W} \in \mathbb{R}^{d \times T}$ where columns correspond to \mathbf{w}_i ,

$$\mathbf{h}_t = \mathbf{VW}\mathbf{a}_t$$

- And finally

$$\mathbf{H} = \mathbf{VWA}$$

- How to compute the matrix \mathbf{A} ?

$$\mathbf{A} = \text{softmax}(\mathbf{W}^\top \mathbf{K}^\top \mathbf{Q} \mathbf{W})$$

where the softmax is applied column-wise.

- Why softmax? to get positive entries, and columns summing to 1.
- Why $\mathbf{W}^\top \mathbf{K}^\top \mathbf{Q} \mathbf{W}$? Bilinear form over the input

Combining vectors with attention

- Putting everything together:

$$\mathbf{H} = \mathbf{V}\mathbf{W}\text{softmax}(\mathbf{W}^T\mathbf{K}^T\mathbf{Q}\mathbf{W})$$

where $\mathbf{H}, \mathbf{W} \in \mathbb{R}^{d \times T}$ and $\mathbf{V}, \mathbf{K}, \mathbf{Q} \in \mathbb{R}^{d \times d}$

- $\mathbf{V}, \mathbf{K}, \mathbf{Q}$ are parameters to be learned.
- This operation is called **self-attention**

Combining vectors with attention

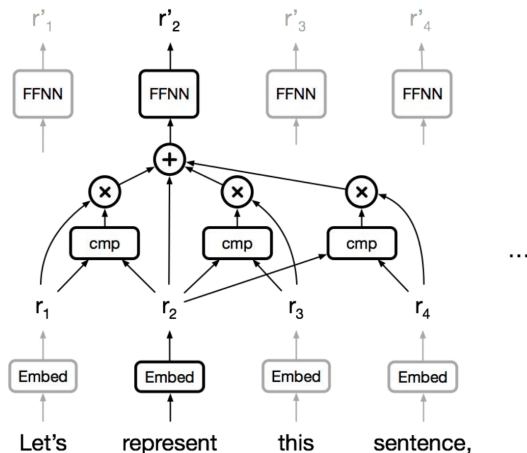
- Putting everything together:

$$\mathbf{H} = \mathbf{V}\mathbf{W}\text{softmax}(\mathbf{W}^\top \mathbf{K}^\top \mathbf{Q}\mathbf{W})$$

where $\mathbf{H}, \mathbf{W} \in \mathbb{R}^{d \times T}$ and $\mathbf{V}, \mathbf{K}, \mathbf{Q} \in \mathbb{R}^{d \times d}$

- $\mathbf{V}, \mathbf{K}, \mathbf{Q}$ are parameters to be learned.
- This operation is called **self-attention**
- It can be generalized to **multiple heads**:
 - Split input vectors into n subvectors of dimension d/n ,
 - Apply self attention (with different $\mathbf{V}, \mathbf{K}, \mathbf{Q}$) over these smaller vectors
 - Concatenate the results to get back d dimensional vectors

Combining vectors with attention



from Vaswani and Huang:

<http://web.stanford.edu/class/cs224n/slides/>

Transformer network

Transformer block:


- Multi-head attention layer with skip connection and normalization
- Followed by feed forward with skip connection and normalization

Skip connection+normalization:

- Given a network block **nn** and input **x**
- The output **y** is computed as

$$\mathbf{y} = \mathbf{norm}(\mathbf{x} + \mathbf{nn}(\mathbf{x}))$$

where **norm** normalize the input



Vaswani et al.
(2017)


Transformer network

Feed forward block

- Two layer network, with ReLU activation

$$\mathbf{y} = \mathbf{W}_2 \text{ReLU}(\mathbf{W}_1 \mathbf{x})$$

- Usually, $\mathbf{W}_1 \in \mathbb{R}^{4d \times d}$ and $\mathbf{W}_2 \in \mathbb{R}^{d \times 4d}$
- i.e. hidden layer of dimension $4d$.



Vaswani et al.
(2017)


- **Limitation:** self attention does not take position into account!
- Indeed, shuffling the input gives the same results
- **Solution:** add position encodings.
- Replace the matrix \mathbf{W} by $\mathbf{W} + \mathbf{E}$, where $\mathbf{E} \in \mathbb{R}^{d \times T}$
- \mathbf{E} can be learned, or defined using sin and cos:

$$e_{2i,j} = \sin\left(\frac{j}{10000^{2i/d}}\right)$$
$$e_{2i+1,j} = \cos\left(\frac{j}{10000^{2i/d}}\right)$$

Transformer network

Transformer network:

- Word embeddings + Position embeddings
- Then N transformer blocks (e.g. $N = 12$)
- Softmax classifier (e.g. for language modeling)



Vaswani et al.
(2017)

Masking for Transformer Language Models

- In transformer, \mathbf{h}_t depends on **all** inputs
- Could not be used as is for language modeling
- Solution: use mask in attention, to only use past

- Reminder:

$$\begin{aligned}\mathbf{H} &= \mathbf{VW} \text{softmax}(\mathbf{W}^\top \mathbf{K}^\top \mathbf{QW}) \\ &= \mathbf{VWA}\end{aligned}$$

Hence, \mathbf{a}_{it} is weight of input i in representation of position t

- We want representation at time t to only depends on $i \leq t$
- We could enforce $\mathbf{a}_{it} = 0$ for $i \geq t$

Masked softmax

- We introduce the masked softmax operator
- Given an input \mathbf{x} and a binary mask \mathbf{m} ,

$$[\text{masked_softmax}(\mathbf{x}, \mathbf{m})]_i = \frac{m_i \exp(x_i)}{\sum_{i=1}^d m_i \exp(x_i)}$$

- Still sums to one, $m_i = 0$ implies $[\text{masked_softmax}(\mathbf{x}, \mathbf{m})]_i = 0$
- Sometimes implemented as:

$$\text{softmax}(\mathbf{x} + \log(\mathbf{m}))$$

- **Beware:** do not learn the mask (e.g. PyTorch: `register_buffer`)

Training of a Transformer

- In practice, transformers are very unstable during training
- If the learning rate is too large, it diverges
- However if the learning rate is too small, it does not learn well

Transformer network for Language Modeling: Results

Model	bpc
LSTM	1.25
Transformer	1.07

- Text8
- Character level language modeling
- bpc = bit per character.

Why are language model useful?

- Standard Language Models are not that useful as such
- For specific-tasks we will see that they can be useful in Lecture 6
- For controlled generation (Machine Translation, Speech to Text, Question Answering...) we need more.
- **How to build a "controllable" text generation system using a language model ?**

- Language Model
- **Conditioned Language Model** : focus on Sequence to Sequence

- Problematically, controllable text generation can be seen as estimating:

$P(w_t | w_1, \dots, w_{t-1}, C)$ where C is a conditioning variable

Sequence to Sequence Architecture

Direct modeling of translation

We have:

a sentence $S = (x_1, \dots, x_m)$ in a **Source** language (e.g. French)

its translation $T = (y_1, \dots, y_n)$ in a **Target** language (e.g. English)

We directly work on the probability of a translation given a source sentence by expressing translation as conditional language modeling:

$$P(T | S) = \prod_{t=1}^n P(y_t | y_{t-1}, \dots, y_1, S)$$

Goal Learn a translation model where T is the most probable sentence given S :

$$T = \underset{T' \text{ in Target language}}{\operatorname{argmax}} P(T' | S)$$

Challenge How to encode the source sentence S ?

Sequence to Sequence: Machine Translation

We want to condition a language model of the target language (e.g English) on a source sentence

- ① Encode source sentences
- ② Generate the target sentence based on the encoded source and a language target language model


Sequence to Sequence: Machine Translation

We want to condition a language model of the target language (e.g English) on a source sentence

- ① Encode source sentences
 - ② Generate the target sentence based on the encoded source and a language target language model
- We have seen that Neural Networks are good Language Models (i.e. can generate proper sentences)
 - We have seen that Neural Networks are good at modelling sequence.
 - We are going to combine two network
 - An **encoder** for encoding source sentences
 - A **decoder** for conditioned language modelling

→ This new architecture is referred to as an *encoder-decoder* or *sequence to sequence model*


Simple approach: sequence-to-sequence (seq2seq)



Equivalent to:

1. Build a representation of the source sentence by taking last hidden layer \mathbf{h}_T of LSTM applied to the source sentence
2. Use this representation as initialization of the hidden variables of LSTM applied to the target sentence

Simple approach: sequence-to-sequence (seq2seq)



- Pro:
 - Very simple to implement:
 - input: Concatenation of source and target sentence.
 - output: target sentence
- Cons:
 - Needs large hidden layer to store everything about source sentence
 - Does not work on very long sentences
 - Same conditioning for the whole target sentence

Architecture

- Encoder and Decoder can be any NN architecture seen so far
- In practice, LSTMs and Transformer are the most efficient (in most cases)

Decoding

- At inference, we can improve performance by using *beam-search* instead of *greedy* decoding

- As any Neural Networks, we train a seq2seq architecture with backpropagation
- Using pairs of source-target aligned sentences we train the model to generate the target language based on source language

Source: *This week we'll continue to try to close a deal to purchase a dairy farm.*

Target: *Cette semaine, nous allons continuer d'essayer de signer un contrat d'achat d'une exploitation laitière.*

Attention Mechanism for sequence to sequence

- To overcome the main encoding issue, the sequential attention on the source sentence improve importantly the performance

- How good is a given machine translation system?
- Hard problem, since many different translations acceptable
- Evaluation metrics:
 - subjective judgments by human evaluators
 - automatic evaluation metrics
 - task-based evaluation (how much post-editing effort? does information come across?)

NB: Evaluating sequence generation model is never easy (subjectivity!)

6

⁶from Philipp Koehn: <http://mt-class.org/jhu/>

Measure n -gram overlap between machine translation output and reference translation

Compute precision for n -grams of size 1 to 4

Add brevity penalty to avoid too short translations

$$\text{BLEU} = \min \left(1, \frac{\text{output_length}}{\text{reference_length}} \right) \left(\prod_{i=1}^4 \frac{C_i}{N_i} \right)^{\frac{1}{4}}$$

where C_i is the number of correct n -gram of size i and N_i is the total number of n -grams in the output of the system

Computed over full corpora, not just a sentence from Philipp Koehn: <http://mt-class.org/jhu/>

- (Abstractive) Summarization
 - Input:** Document
 - Output:** Summary
- Text Simplification
 - Input:** Complex sentence
 - Output:** Simplified sentence
- Multi-Modal tasks
 - Speech To Text
 - Caption Generation (Image to Text)

- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., and Sutskever, I. (2019). Language models are unsupervised multitask learners. *OpenAI Blog*, 1(8):9.
- Rae, J. W., Dyer, C., Dayan, P., and Lillicrap, T. P. (2018). Fast parametric learning with activation memorization. *arXiv preprint arXiv:1803.10049*.
- Siegelmann, H. T. and Sontag, E. D. (1992). On the computational power of neural nets. In *Proceedings of the fifth annual workshop on Computational learning theory*, pages 440–449.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. (2017). Attention is all you need. In *Advances in neural information processing systems*, pages 5998–6008.